

CS 140 Project 1: The Rotating Staircase Deadline Scheduler Documentation

Members:

Bugaoisan, Denver Earl Paul S.

Nicolas, Jack Vincent

Ragunton, Carl David B.

Gdrive link for the video documentation:

https://drive.google.com/file/d/1TCENNM3x2FPPSFAOd5CRaKpITJ4kf9gw/view?usp=share_link

1. [1,2,3,4,5] List of all phases with working code (i.e., which branches are eligible for checking)

The branches that our group was able to do are the following:

- phase1
- phase2
- phase3
- phase4
- phase5

2. [1,2,3,4,5] All references used with purpose specified (if any; otherwise, explicitly state that you did not use any external resource)

The only reference that our group used for this project is the “CS 140 Laboratory 5: Process Scheduling” from

https://uvle.upd.edu.ph/pluginfile.php/767851/mod_resource/content/4/CS_140_22_1_LAB_5.pdf. The purpose of using this reference is to gain ideas and to replicate how the Round Robin Scheduler works in the lab. That will serve as the base of how the whole project works. The syscall schedlog is also mainly from this activity.

3. [1,2,3,4,5] List of all global variables introduced with purpose specified (only those not related to system calls)

```
int num in proc.c
```

The purpose of this global variable is to assign an ascending number to all processes. Once a process is made, the num increments and it is assigned to the process. The same idea applies when a process is used. This way, we can know which process is more recently enqueued and which ones are there first. This variable ensures the application of FIFO in the scheduler.

```
int level_quantum_active [RSDL_LEVELS];  
int level_quantum_expired [RSDL_LEVELS]; in proc.c
```

The purpose of these global arrays is to store the remaining quantum per level. The level_quantum_active is responsible for the levels in the active set, while the level_quantum_expired is responsible for the levels in the expired set.

4. [1,2,3,4,5] List of changes made to the process control block

There are several variables that were added in the PCB:

```
int ticks_left;
int queue_num;
int active;
int level;
int quantum_level;
int reset;
int prio;
```

Variables	Description
<code>int ticks_left;</code>	The quantum of the process
<code>int queue_num;</code>	The queue number in local-level
<code>int active;</code>	Whether the process is in the active set (<code>active = 1</code>) or expired set (<code>active = 0</code>)
<code>int level;</code>	The priority level of the process
<code>int quantum_level;</code>	This determines whether the quantum for the current level is depleted.
<code>int reset;</code>	This dictates whether the quanta per level are set to <code>RSDL_LEVEL_QUANTUM</code> (<code>reset = 1</code>). Otherwise, continue.
<code>int prio;</code>	Starting level if <code>priofork</code> syscall is called.

5. [1,2,3,4,5] Code representation of the active set and its levels.

First, it is important to take note of the PCB. The highlighted line shows the int variable active. This determines whether the process is in the active set (when active = 1) or not (when active = 0). This is in proc.h.

```
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // switch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
    char name[16];          // Process name (debugging)
    int ticks_left;
    int queue_num;
    int active;
    int level;
    int quantum_level;
    int reset;
    int prio;
};
```

This is the code for the active set and its levels. This is located in proc.c, under the scheduler function.

```
for (int l = 0; l < RSDL_LEVELS; l++){
    cprintf("%d|active|%(%d)", ticks, l, level_quantum_active[l]);

    struct proc *pp;

    char* initname="init";
    for (int k = 0; k < NPROC; k++) {
        pp = &ptable.proc[sorted[k]];
        if (pp->level != l) continue;
        if (pp->pid == 0 && pp->name != initname) continue;
        if (pp->state == 5 || pp->state == EMBRYO || pp->state == UNUSED) continue;

        if (pp->state == RUNNING && pp->active == 1) cprintf("[%d]s:%d(%d)", pp->pid, pp->name, pp->state, pp->ticks_left);
        else if (pp->active == 1) cprintf("[%d]s:%d(%d)", pp->pid, pp->name, pp->state, pp->ticks_left);
    }
    cprintf("\n");
}
```

The starting for loop is responsible for the different levels that will be printed out in the active set. This is followed by a `cprintf` which prints out the tick, level and quantum in that level.

The levels are checked whether the processes are supposed to be printed on the current level or not, and if so, we just skip and loop again. The pid of the processes are also checked to skip when it is 0 and the name is not equal to the `initname`, which means that they are the killed processes.

The nested for loop is used to iterate through the active processes (when `{process} -> active = 1`) and print them out for the current level as long as their `proc_state` is not `UNUSED`, `EMBRYO` or `ZOMBIE`.

```
146|active|0(100),[1]init:2(9),[2]sh:2(10),[3]test:2(9)
146|active|1(100)
146|active|2(100),[4]test:4(10)
146|active|3(100),[5]test:2(9)
146|active|4(100),[6]test:2(9)
```

6. [1,2,3,4,5] Implementation explanation for initial enqueueing of new processes.

Our group did not use literal queues to create a queue. Instead, we use values and conditions to replicate the queue structure.

```
37 // Per-process state
38 struct proc {
39     uint sz; // Size of process memory (bytes)
40     pde_t* pgdir; // Page table
41     char *kstack; // Bottom of kernel stack for this process
42     enum procstate state; // Process state
43     int pid; // Process ID
44     struct proc *parent; // Parent process
45     struct trapframe *tf; // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan; // If non-zero, sleeping on chan
48     int killed; // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd; // Current directory
51     char name[16]; // Process name (debugging)
52     int ticks_left;
53     int queue_num;
```

We added queue_num to the PCB. This will represent when a process “enters the imaginary queue.” The higher the queue_num, the more recent a process enters the queue. This ensures the FIFO.

```
9 int num = 0; //phase1
```

In proc.c, we added a global variable. This will give the queue_num of the processes their values.

```
91 found:
92     p->state = EMBRYO;
93     p->pid = nextpid++;
94     p->queue_num = num; //phase1
95     num += 1;
96     p->active = 1; //phase2
97     p->ticks_left = RSDL_PROC_QUANTUM;
98     p->level = RSDL_STARTING_LEVEL;
```

This is where the initial enqueueing of processes happens. When a process is created, it's queue->num becomes num. Then, num is incremented to let the next process have a higher value. That means the lower the queue_num, the deeper it is in the queue.

For multiple levels, we just set conditions on which level will be used.

7. [1,2,3,4,5] Implementation explanation for dequeuing of exiting processes.

First, we look into the proc.h and find the PCB with a variable queue_num. This represents when a process “enters the imaginary queue.” The higher the queue_num, the more recent a process enters the queue. This ensures the FIFO.

```
// Per-process state
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // switch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
    char name[16];          // Process name (debugging)
    int ticks_left;
    int queue_num;
    int active;
    int level;
    int quantum_level;
    int reset;
    int prio;
};
```

Meanwhile, in proc.c, we have a global variable called num. This will give the queue_num of the processes their values.

```
int num = 0;
```

Whenever, these pairs come up in the code, it dequeues the exiting process. When num is incremented it enables the next process to have a higher value. This is an instantaneous dequeue and enqueue. As the value of queue_num goes up, the process is dequeued from its original position and enqueued to the end of the “queue”.

```
num+=1;
p->queue_num = num;
```


8. [1,2,3,4,5] Implementation explanation for process-local quantum consumption.

First, the process-local quantum is initialized inside the for (;;) in the scheduler function under proc.c. Each process is loaded with the process quantum of RSDL_PROC_QUANTUM.

```
if (p->active == 1) p->ticks_left = RSDL_PROC_QUANTUM;
```

Now, the main code responsible for the quantum consumption in the local level is located in trap.c. In the trap function, when the process is running, the ticks_left is decremented for each tick that it runs. When, the ticks_left == 0 (this means that the process-local quantum is fully depleted), then we yield.

```
if(myproc() && myproc()->state == RUNNING &&
    tf->trapno == T_IRQ0+IRQ_TIMER){
    --myproc()->quantum_level;
    --myproc()->ticks_left;
    if (myproc()->quantum_level <= 0 || myproc()->ticks_left == 0){
        yield();
    }
}
```

9. [1,2,3,4,5] Implementation explanation for process-local quantum replenishment.

First, the process-local quantum is initialized inside the for (;;) in the scheduler function under proc.c. Each process is loaded with the process quantum of RSDL_PROC_QUANTUM.

```
if (p->active == 1) p->ticks_left = RSDL_PROC_QUANTUM;
```

The replenishment of the process-local quantum can be found in proc.c under the scheduler function. It mainly occurs when the ticks_left of the process has been depleted. It is replenished with a value depending on the RSDL_PROC_QUANTUM.

```
else if (p->ticks_left == 0 && p->active == 1){  
    num+=1;  
    p->queue_num = num;  
    p->level++;  
    p->ticks_left=RSDL_PROC_QUANTUM;  
    if (p->level == RSDL_LEVELS) {  
        if (p->prio != -1) p->level = p->prio;  
        else p->level=RSDL_STARTING_LEVEL;  
        p->active=0;  
        p->ticks_left=RSDL_PROC_QUANTUM;  
    }  
}
```

As for replenishment of the process-local quanta in the expired set, they are replenished when they are moved from the active set to the expired set.

```
if (p->level == RSDL_LEVELS) {  
    if (p->prio != -1) p->level = p->prio;  
    else p->level=RSDL_STARTING_LEVEL;  
    p->active=0;  
    p->ticks_left=RSDL_PROC_QUANTUM;  
}
```

10. [1,2,3,4,5] Implementation explanation for schedlog.

The `schedlog` works in a similar fashion to the `schedlog` in Lab 5. As stated in Lab 5, `schedlog` takes in an integer that corresponds to the length of time (measured in ticks) to output scheduling information. The scheduler was modified to print out the process table before it switches to a new scheduled process if `schedlog` is enabled, or disable `schedlog` if it has reached the specified number of ticks.

```
324  int schedlog_active = 0;
325  int schedlog_lasttick = 0;
326
327  void schedlog(int n) {
328      schedlog_active = 1;
329      schedlog_lasttick = ticks + n;
330  }
```

If the `schedlog` is active, the process table will be first sorted based on their queue number (`queue_num`). The PCB of each process has its own queue number. The sorting algorithm is implemented this way:

Queue numbers of each process is stored in the integer array called "array". There is also an integer array called "sorted" that is used to store the indices of the elements in "array" after they are sorted. The outer for loop with the variable "k" iterates through each element in the "array" and "sorted" arrays. The inner for loop with the variable "j" compares the element at index "k" in the "array" to each element with an index greater than "k". If the element at index "k" is greater than the element at index "j", the elements at these indices are swapped in both the "array" and "sorted" arrays. This process is repeated until all of the elements in the array are in the desired order. After the nested for loops complete, the "array" will be sorted in ascending order and the "sorted" array will contain the indices of the elements in the original "array" after they were sorted.

```

463         int sorted[NPROC];
464         int array[NPROC];
465         for (int k = 0; k < NPROC; k++) {
466             sorted[k]=k;
467             struct proc *pp;
468             pp = &ptable.proc[k];
469             array[k]=pp->queue_num;
470         }
471
472         for (int k = 0; k < NPROC; k++) {
473             for (int j = k+1; j < NPROC; j++) {
474                 if(array[k] > array[j]) {
475                     int temp = sorted[k];
476                     sorted[k]=sorted[j];
477                     sorted[j]=temp;
478
479
480                     int temp2 = array[k];
481                     array[k] = array[j];
482                     array[j] = temp2;
483                 }
484             }
485         }

```

In order to print the scheduling information:

The priority levels are represented by the integer value of the "level" field of each process. The outer for loop with the variable "l" iterates through each priority level. It prints first the global ticks, set (active or expired), levels and the quantum of each level. The inner for loop with the variable "k" iterates through each process in the priority queue. For each process, the code checks if the process is at the current priority level, and if the process is in a valid state (not EMBRYO, UNUSED, or ZOMBIE). If these conditions are met, another if condition is used to check if the state of the process is running and active = 1. If true, the code prints out information about the process, including its PID, name, state, and the number of ticks remaining. Printing for the expired set works in a similar fashion with active = 0.

```
488     for (int l = 0; l < RSDL_LEVELS; l++){
489         cprintf("%d|active|%d(%d)", ticks, l, level_quantum_active[l]);
490
491         struct proc *pp;
492
493         char* initname="init";
494         for (int k = 0; k < NPROC; k++) {
495             pp = &ptable.proc[sorted[k]];
496             if (pp->level != l) continue;
497             if (pp->pid == 0 && pp->name != initname) continue;
498             if (pp->state == 5 || pp->state == EMBRYO || pp->state == UNUSED) continue;
499
500             if (pp->state == RUNNING && pp->active == 1) cprintf(", [%d]%s:%d(%d)", pp->pid,
501                 pp->name, pp->state, pp->ticks_left);
502             else if (pp->active == 1) cprintf(", [%d]%s:%d(%d)", pp->pid, pp->name, pp->state,
503                 pp->ticks_left);
504         }
505         cprintf("\n");
506     }
```

11. [..2,3,4,5] Code representation of the expired set and its levels.

First, it is important to take note of the PCB. The highlighted line shows the int variable active. This determines whether the process is in the expired set (when active = 0) or not (when active = 1). This is in proc.h.

```
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // switch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
    char name[16];          // Process name (debugging)
    int ticks_left;
    int queue_num;
    int active;
    int level;
    int quantum_level;
    int reset;
    int prio;
};
```

This is the code for the expired set and its levels. This is located in proc.c, under the scheduler function.

```
for (int l = 0; l < RSDL_LEVELS; l++){
    cprintf("%d|expired|%d(%d)", ticks,l, level_quantum_expired[l]);
    struct proc *pp;

    char* initname="init";
    for (int k = 0; k < NPROC; k++) {
        pp = &ptable.proc[sorted[k]];
        if (pp->level != l) continue;
        if (pp->pid == 0 && pp->name != initname) continue;
        if (pp->state == 5 || pp->state == EMBRYO || pp->state == UNUSED) continue;

        if (pp->state == RUNNING && pp->active == 0) cprintf("[%d]s:%d(%d)", pp->pid, pp->name, pp->state, pp->ticks_left);
        else if (pp->active == 0) cprintf("[%d]s:%d(%d)", pp->pid, pp->name, pp->state, pp->ticks_left);
    }
    cprintf("\n");
}
```

The starting for loop is responsible for the different levels that will be printed out in the expired set. This is followed by a cprintf which prints out the tick, level and quantum in that level.

The levels are checked whether the processes are supposed to be printed on the current level or not, and if so, we just skip and loop again. The pid of the processes are also checked to skip when it is 0 and the name is not equal to the initname, which means that they are the killed processes.

The nested for loop is used to iterate through the expired processes (when {process}->active = 0) and print them out for the current level as long as their proc_state is not UNUSED, EMBRYO or ZOMBIE.

```
146|expired|0(100)
146|expired|1(100)
146|expired|2(100)
146|expired|3(100)
146|expired|4(100)
```

12. [..2,3,4,5] Implementation explanation for transferring a process from the active to expired

The PCB of each process has `ticks_left` and `active` variables as shown below.

`ticks_left` acts as the quantum for the process and `active` acts if the process is in the active set or expired set.

```
38 struct proc {
39     uint sz;                // Size of process memory (bytes)
40     pde_t* pgdir;          // Page table
41     char *kstack;           // Bottom of kernel stack for this process
42     enum procstate state;   // Process state
43     int pid;                // Process ID
44     struct proc *parent;    // Parent process
45     struct trapframe *tf;   // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan;             // If non-zero, sleeping on chan
48     int killed;             // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd;      // Current directory
51     char name[16];          // Process name (debugging)
52     int ticks_left;         // OozoraCielo, 2 weeks ago * temp commit
53     int queue_num;
54     int active;             // OozoraCielo, 2 days ago * temp commit
55     int level;
56     int quantum_level;
57     int reset;
58     int prio;
```

If the process is still active (`active = 1`); quantum of the process is depleted (`p->ticks_left = 0`); and it is in the last level of the active set (`p->level=RSDL_LEVELS`), the process will be transferred to the expired set by setting `active=0` as shown in the code below.

```
466     if (p->ticks_left == 0 && p->active == 1){
467         num+=1;
468         p->queue_num = num;
469         p->level++;
470         p->ticks_left=RSDL_PROC_QUANTUM;
471     };
472     if (p->level == RSDL_LEVELS) {
473         p->level=RSDL_STARTING_LEVEL;
474         p->active=0;
475         p->ticks_left=RSDL_PROC_QUANTUM;
476     }
```


13. [..2,3,4,5] Implementation explanation for swapping of sets.

The scheduler will check if all the runnable processes are still active. As shown in the code below, if the variable `all_expired` is equal to 0, a runnable process is still active (`ppp->active = 1`). If all the processes are not active, `all_expired` is equal to 1.

```
590     int all_expired = 1;
591     for (int k = 0; k < NPROC; k++) {
592         ppp = &ptable.proc[k];
593         if (ppp->level < currlevel) continue;
594
595         if (ppp->state != RUNNABLE) continue;
596         else if (ppp->active == 1){
597             all_expired = 0;
598         }
599     }
```

If `all_expired` is equal to 1, `reset` will toggle and the processes are arranged in FIFO manner. A loop is also used to set the unrunnable processes to 0, set their quantum to `RSDL_PROC_QUANTUM` and arrange them in FIFO manner.

```
601     if (all_expired == 1){
602         p->reset = 1;
603         num += 1; //phase1F
604         p->queue_num = num; //phase1
605
606         for (int k = 0; k < NPROC; k++) {
607             ppp = &ptable.proc[k];
608             if (ppp->level != currlevel) continue;
609             if (ppp->active == 1){
610                 ppp->active = 0;
611                 ppp->ticks_left = RSDL_PROC_QUANTUM;
612                 num += 1; //phase1F
613                 ppp->queue_num = num; //phase1
614             }
615         }
```

Lastly, since all processes now are not active, they are set to active again by another for loop.

```
616         for (int k = 0; k < NPROC; k++) {
617             ppp = &ptable.proc[k];
618             ppp->active = 1;
619         }
620     }
```

14. [...3,4,5] Implementation explanation for downgrading of process levels

The PCB of each process has a variable level.

```
38 struct proc {
39     uint sz;                // Size of process memory (bytes)
40     pde_t* pgdir;          // Page table
41     char *kstack;          // Bottom of kernel stack for this process
42     enum procstate state;   // Process state
43     int pid;               // Process ID
44     struct proc *parent;    // Parent process
45     struct trapframe *tf;   // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan;            // If non-zero, sleeping on chan
48     int killed;            // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd;     // Current directory
51     char name[16];         // Process name (debugging)
52     int ticks_left;
53     int queue_num;
54     int active;
55     int level;              // OozoraCielo, yesterday * temp commit ...
56     int quantum_level;
57     int reset;
58     int prio;
59 };
```

Downgrading the level of each process is based on the `ticks_left` of the certain process; `ticks_left` also serves as the quantum for each process. If `ticks_left` is equal to 0 or the quantum is depleted, the priority level of the process will downgrade through `p->level++` in a FIFO manner and its quantum will also be replenished. If the level of the process is equal to `RSDL_LEVELS`, the process will be sent to the expired set based on `RSDL_STARTING_LEVEL`.

```
558     if (p->quantum_level<=0){
559         for (int k = 0; k < NPROC; k++){
560             struct proc *ppp;
561             ppp = &ptable.proc[sorted[k]];
562             if (ppp->level != currlvl) continue;
563             num++;
564             ppp->queue_num = num;
565             ppp->level += 1;
566             ppp->ticks_left=RSDL_PROC_QUANTUM;
567             if (ppp->level >= RSDL_LEVELS) {
568                 if (ppp->prio != -1) ppp->level = ppp->prio;
569                 else ppp->level=RSDL_STARTING_LEVEL;
570                 ppp->active=0;
571                 ppp->ticks_left=RSDL_PROC_QUANTUM;
572             }
573         }
574         num++;
575         p->queue_num = num;
576     }
577     else if (p->ticks_left == 0 && p->active == 1){
578         num++;
579         p->queue_num = num;
580         p->level++;
581         p->ticks_left=RSDL_PROC_QUANTUM;
582         if (p->level == RSDL_LEVELS) {
583             if (p->prio != -1) p->level = p->prio;
584             else p->level=RSDL_STARTING_LEVEL;
585             p->active=0;
586             p->ticks_left=RSDL_PROC_QUANTUM;
587         }
588     }
```

15. [.....4,5] Implementation explanation for level-local quantum consumption

To implement this part, 2 changes are made to the PCB.

```
37 // Per-process state
38 struct proc {
39     uint sz; // Size of process memory (bytes)
40     pde_t* pgdir; // Page table
41     char *kstack; // Bottom of kernel stack for this process
42     enum procstate state; // Process state
43     int pid; // Process ID
44     struct proc *parent; // Parent process
45     struct trapframe *tf; // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan; // If non-zero, sleeping on chan
48     int killed; // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd; // Current directory
51     char name[16]; // Process name (debugging)
52     int ticks_left;
53     int queue_num;
54     int active;
55     int level;
56     int quantum_level;
57     int reset;
58     int prio;
59 };
```

They are the `quantum_level` and `reset`. `quantum_level` will let the program know what is the quantum of the current level of the process. `reset` will signal if the quanta of the levels should be reset.

In order to make the decrement of ticks of the `quantum_level` similar to the process of the `ticks_left`, the `trap.c` has also been changed.

```
103 // Force process to give up CPU on clock tick.
104 // If interrupts were on while locks held, would need to check nlock.
105 if(myproc() && myproc()->state == RUNNING &&
106     tf->trapno == T_IRQ0+IRQ_TIMER){
107     --myproc()->quantum_level;
108     --myproc()->ticks_left;
109     if (myproc()->quantum_level <= 0 || myproc()->ticks_left == 0){
110         yield();
111     }}
```

This causes the quantum level to decrement every tick just like how `ticks_left` works. Also, on line 109, I included an OR statement inside the if condition. Thus, whenever one of the `quantum_level` or the `ticks_left` are 0, the process will be forced to yield the control of the cpu.

On the `proc.c`, 2 global arrays has been added:

```

10 int level_quantum_active [RSDL_LEVELS]; // phase 4
11 int level_quantum_expired [RSDL_LEVELS];

```

As their name suggests, level_quantum_active will contain all the quantum for active level sets and level_quantum_expired is for the quantum of expired level sets.

```

393 for (int i = 0; i < RSDL_LEVELS; i++) {
394     level_quantum_active[i] = RSDL_LEVEL_QUANTUM;
395     level_quantum_expired[i] = RSDL_LEVEL_QUANTUM;
396 }
397 for (int k = 0; k < NPROC; k++) {
398     struct proc *ppp;
399     ppp = &ptable.proc[k];
400     if (ppp->active == 1) {
401         ppp->quantum_level = level_quantum_active[ppp->level];
402     }
403     else {
404         ppp->quantum_level = level_quantum_expired[ppp->level];
405     }
406 }

```

The first for loop here initializes the value of all the elements of the 2 arrays into the value of RSDL_LEVEL_QUANTUM. For the second loop, it assigns the quantum_level of each process to the correct one, depending if it is active or expired.

```

486 if (p->quantum_level < 0) p->quantum_level=0;
487 level_quantum_active[p->level]=p->quantum_level;

```

Inside the code for schedlog, it is made sure that the values never go down zero. The level_quantum_active is also updated here to let the other processes later to be able to access their correct quantum_level.

```

488 for (int l = 0; l < RSDL_LEVELS; l++) {
489     cprintf("%d|active|%d(%d)", ticks, l, level_quantum_active[l]);
490 }
507 for (int l = 0; l < RSDL_LEVELS; l++) {
508     cprintf("%d|expired|%d(%d)", ticks, l, level_quantum_expired[l]);
509 }

```

The part of the code that prints the quantum of the level is above separated for the active and expired sets. They just use l since that is their correct level.

```

529     int sorted[NPROC];
530     int array[NPROC];
531     for (int k = 0; k < NPROC; k++) {
532         sorted[k]=k;
533         struct proc *pp;
534         pp = &ptable.proc[k];
535         array[k]=pp->queue_num;
536     }
537
538
539     for (int k = 0; k < NPROC; k++) {
540         for (int j = k+1; j < NPROC; j++) {
541             if(array[k] > array[j]) {
542                 int temp = sorted[k];
543                 sorted[k]=sorted[j];
544                 sorted[j]=temp;
545
546
547                 int temp2 = array[k];
548                 array[k] = array[j];
549                 array[j] = temp2;
550             }
551         }
552     }
553

```

After the part for schedlog, we do a imaginary sort again to be able to use the correct index later.

```

554     int currlevel = p->level;
555     if (p->quantum_level < 0) p->quantum_level=0;
556     level_quantum_active[currlevel]=p->quantum_level;
557

```

Here, we just make sure that the level quantum will never go negative and update the level_quantum_active again.

```

554     int currlevel = p->level;
555     if (p->quantum_level < 0) p->quantum_level=0;
556     level_quantum_active[currlevel]=p->quantum_level;
557
558     if (p->quantum_level<=0){
559         for (int k = 0; k < NPROC; k++){
560             struct proc *ppp;
561             ppp = &ptable.proc[sorted[k]];
562             if (ppp->level != currlevel) continue;
563             num+=1;
564             ppp->queue_num = num;
565             ppp->level += 1;
566             ppp->ticks_left=RSDL_PROC_QUANTUM;
567             if (ppp->level >= RSDL_LEVELS) {
568                 if (ppp->prio != -1) ppp->level = ppp->prio;
569                 else ppp->level=RSDL_STARTING_LEVEL;
570                 ppp->active=0;
571                 ppp->ticks_left=RSDL_PROC_QUANTUM;
572             }
573         }
574         num+=1;
575         ppp->queue_num = num;
576     }
577     else if (p->ticks_left == 0 && p->active == 1){
578         num+=1;
579         p->queue_num = num;
580         p->level++;
581         p->ticks_left=RSDL_PROC_QUANTUM;
582         if (p->level == RSDL_LEVELS) {
583             if (p->prio != -1) p->level = p->prio;
584             else p->level=RSDL_STARTING_LEVEL;
585             p->active=0;
586             p->ticks_left=RSDL_PROC_QUANTUM;
587         }
588     }

```

From the last phase, this part is just a single if statement. But we can't combine the scenario of the quantum_level being zero and the ticks_left being zero into one because the p->level of the current process will go up twice. Thus, we first check for the quantum_level having a value of 0. If it is, all processes in the current level will go to the next level. Since, the queue here is just a visualization, we can just additionally increase the queue_num of the current process to let it be the last process to be enqueued.

The else if statement below is just the same from the last phase.

```

601         if (all_expired == 1){
602             p->reset = 1;
603             num += 1; //phase1F
604             p->queue_num = num; //phase1
605

```

Here is when the reset of the PCB comes in. If all the runnable processes are in the expired set, we need to reset the quantum of all the levels because a swap will be happening.

```

620         if (p->reset == 1){
621             for (int i = 0; i < RSDL_LEVELS; i++){
622                 level_quantum_active[i] = RSDL_LEVEL_QUANTUM;
623             }
624             p->reset = 0;
625         }
626

```

If a swap is to be done, a reset should also be done. We just initialize the quantum of all the levels of active sets to be RSDL_LEVEL_QUANTUM again. Then, the reset should stay 0 again till another reset will be done.

16. [.....5] Implementation explanation for priofork

To implement priofork, we changed the required files to create a syscall.
user.h

```
30  int priofork(int);
```

usys.S

```
35  SYSCALL(priofork)
```

syscall.h

```
26  #define SYS_priofork 25
```

Syscall.c

```
109  extern int sys_priofork(void);
```

```
136  [SYS_priofork] sys_priofork,
```

Sysproc.c

```
16  int
17  sys_priofork(void)
18  {
19      int priority;
20
21      if(argint(0, &priority) < 0)
22          return -1;
23      return priofork(priority);
24  }
```

This will pass the value of the priofork called into the function of priofork.

Defs.h

```
125  int      priofork(int);
```

Aside from those, there is also a change in the PCB in proc.h.

```
58  int prio;
```

int prio has been added to it. This will tell us later whether fork or priofork has been used to create a process alongside with the value used by priofork.

The remaining changes are in the proc.c.

Our implementation of priofork is just the same as fork. We just duplicated the fork function in proc.c to create a function named priofork. The difference is that priofork has an int as a parameter.

```
226     np->state = RUNNABLE;  
227     np->prio = -1;  
228
```

This is in the fork function. We just set np->prio to -1 which will be passed in the ptable. If prio is -1, it means that the process is created by fork.

```
234 int
235 priofork(int priority)
236 {
237     int i, pid;
238     struct proc *np;
239     struct proc *curproc = myproc();
240
241     // Allocate process.
242     if((np = allocproc()) == 0){
243         return -1;
244     }
245
246     // Copy process state from curproc.
247     if((np->pgdir = copyuvm(curproc->pgdir, curproc->sz)) == 0){
248         kfree(np->kstack);
249         np->kstack = 0;
250         np->state = UNUSED;
251         return -1;
252     }
253     np->sz = curproc->sz;
254     np->parent = curproc;
255     *np->tf = *curproc->tf;
256
257     // Clear %eax so that fork returns 0 in the child.
258     np->tf->eax = 0;
259
260     for(i = 0; i < NOFILE; i++)
261         if(curproc->ofile[i])
262             np->ofile[i] = filedup(curproc->ofile[i]);
263     np->cwd = idup(curproc->cwd);
264
265     safestrcpy(np->name, curproc->name, sizeof(curproc->name));
266
267     pid = np->pid;
268 }
```

```

268
269     acquire(&ptable.lock);
270
271     np->state = RUNNABLE;
272     np->prio = priority;
273     np->level = priority;
274
275     release(&ptable.lock);
276
277     return pid;
278 }

```

This is the whole priofork function. It only has small changes from fork.

```

271     np->state = RUNNABLE;
272     np->prio = priority;
273     np->level = priority;
274

```

When priofork is called, np->prio will become the value used by priofork. We also set the level of the process to it. This will override the previous level at this point which is the RSDL_STARTING_LEVEL.

```

558     if (p->quantum_level<=0){
559         for (int k = 0; k < NPROC; k++){
560             struct proc *ppp;
561             ppp = &ptable.proc[sorted[k]];
562             if (ppp->level != currlevel) continue;
563             num+=1;
564             ppp->queue_num = num;
565             ppp->level += 1;
566             ppp->ticks_left=RSDL_PROC_QUANTUM;
567             if (ppp->level >= RSDL_LEVELS) {
568                 if (ppp->prio != -1) ppp->level = ppp->prio;
569                 else ppp->level=RSDL_STARTING_LEVEL;
570                 ppp->active=0;
571                 ppp->ticks_left=RSDL_PROC_QUANTUM;
572             }
573         }
574         num+=1;
575         ppp->queue_num = num;
576     }
577     else if (p->ticks_left == 0 && p->active == 1){
578         num+=1;
579         p->queue_num = num;
580         p->level++;
581         p->ticks_left=RSDL_PROC_QUANTUM;
582         if (p->level == RSDL_LEVELS) {
583             if (p->prio != -1) p->level = p->prio;
584             else p->level=RSDL_STARTING_LEVEL;
585             p->active=0;
586             p->ticks_left=RSDL_PROC_QUANTUM;
587         }
588     }
589

```

The last changes are here. When a process is to exceed the max level, the prio is checked whether it is created by for or priofork. If it is created by priofork, instead of resetting its level to the RSDL_STARTING_LEVEL, we set it to its prio.